

# Introduction to Operating Systems

Through tracing, analysis, and experimentation

---

George V. Neville-Neil

February 2017

# Communication

---

# Inter-Process Communication

---

# Goals of IPC

- Data Sharing
- Signaling events
- Control of multiple processes

- Shared files
- Semaphores and Mutexes
- Signals
- Sockets

# Relationship to Networking

- Extension of mechanisms across machines
- Everything is a byte stream
- No record boundaries
- File like API

- Based on hardware interrupt model
- Not useful for data transfer
- Catch and process

# Signal Handling

- Source raises a signal
- destination catches the signal
- Uncaught signals cause a program to exit



# Available Signals

Name	Meaning	Name	Meaning
SIGHUP	line hangup	SIGURG	urgent condition present on socket
SIGINT	interrupt program	SIGSTOP	stop (cannot be caught or ignored)
SIGQUIT	quit program	SIGTSTP	stop signal generated from keyboard
SIGILL	illegal instruction	SIGCONT	continue after stop
SIGTRAP	trace trap	SIGCHLD	child status has changed
SIGABRT	abort program	SIGTTIN	background read attempted from control terminal
SIGEMT	emulate instruction executed	SIGTTOU	background write attempted to control terminal
SIGFPE	floating-point exception	SIGIO	I/O is possible on a descriptor
SIGKILL	kill program	SIGXCPU	cpu time limit exceeded
SIGBUS	bus error	SIGXFSZ	file size limit exceeded
SIGSEGV	segmentation violation	SIGVTALRM	virtual time alarm
SIGSYS	non-existent system call invoked	SIGPROF	profiling timer alarm
SIGPIPE	write on a pipe with no reader	SIGWINCH	Window size change
SIGALRM	real-time timer expired	SIGINFO	status request from keyboard
SIGTERM	software termination signal	SIGUSR1	User defined signal 1
		SIGUSR2	User defined signal 2
		SIGTHR	thread interrupt
		SIGLIBRT	real-time library interrupt

# Tracing Signals

- Earliest bulk data IPC
- Key innovation of UNIX systems
- Depends on file descriptors
  - *STDIN, STDOUT, STDERR*

# Pipe Demonstration

# Internetworked Communication

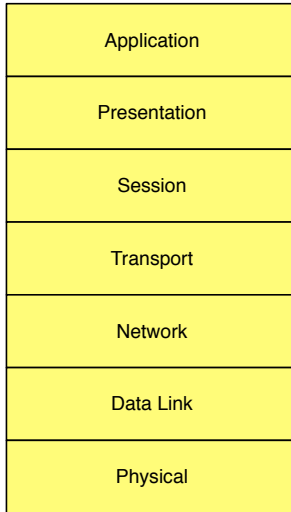
---

- Everyone's TCP/IP Stack
- IPv4, IPv6, UDP, TCP, SCTP
- Various drivers
- Multiple firewalls

# Networking: The ISO Model

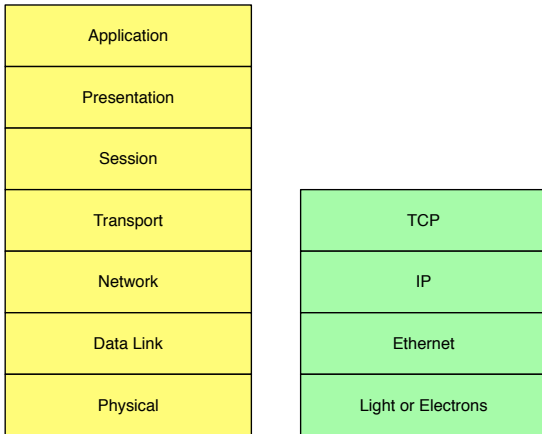
- Canonical description of network protocols
- Each protocols are layered
- Seven layers in all
- Beware Van Jacobsen's warning!

# Networking and Layering






# Networking and Layering



# Networking and Layering



# The User Program View

- User programs use sockets
- Network programs follow UNIX model
- Flexible interfaces for different protocols

- Main programmer interface to networking
- Generic API
- Attempts to support read/write semantics

# Socket System Calls

**socket** Returns a file descriptor

**connect** Connect to a remote program

**bind** Bind a socket to a port

**listen** Listen for connections

**accept** Returns a new file descriptor

## Transferring Data on Sockets

**read** Just like a file

**write** Just like a file

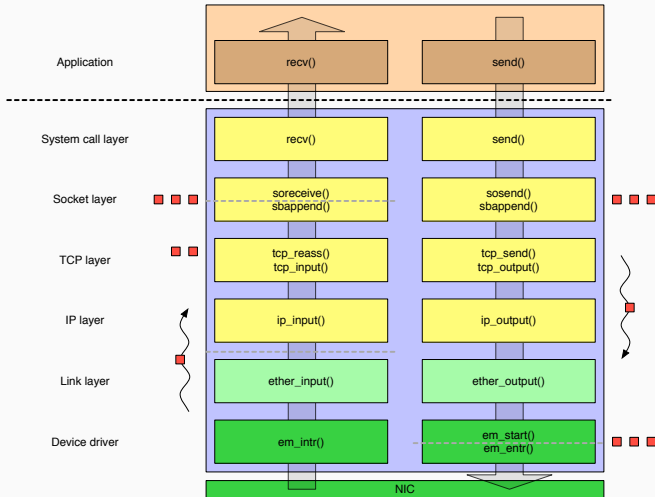
**recv** Receive a single message

**send** Send a single message

**recvmsg** Receive a message with meta-data

**sendmsg** Send a message with meta-data

# Network Stack Overview



- Simplest transport protocol
- No states to maintain
- Data is sent immediately
- Supports multicast
- Only probes are `send` and `receive`




- Transmission Control Protocol
- Stream based
- In order delivery
- Maintains the illusion of a byte stream

# Three Way Handshake

- Initiating a connection between two nodes
  1. Start a connection with a Synchronize (SYN) packet.
  2. Acknowledge the first SYN and initiate a full connection (SYN/ACK)
  3. Acknowledge the second SYN.

# Starting a Connection



**CLOSED**

**SYN SENT** Client initiated a connection.

**SYN RECEIVED** Server received initiation from client.

**ESTABLISHED** Client and server can communicate.

**FIN WAIT 1**

**FIN WAIT 2**

**TIME WAIT**


**CLOSE WAIT**

**LAST ACK** Awaiting client's final acknowledgment


- Sequence Numbers
- Acknowledgements
- The sliding window
- Congestion Control

# The Sliding Window

# The Sliding Window




# The Sliding Window





# The Sliding Window



# The Sliding Window



# The Sliding Window



# The Sliding Window



# The Sliding Window

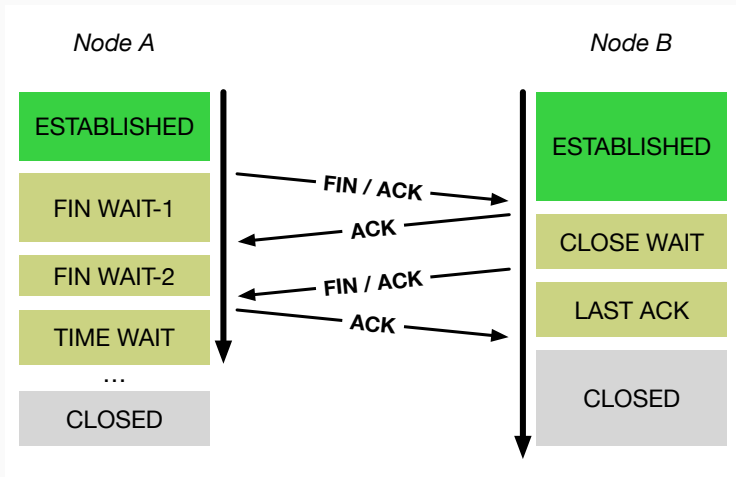


## Four Way Close

- Closing a connection between two nodes
- Each node must close its side of the connection
- More complicated than opening a connection.

1. Node A sends a Finalize (FIN) packet
2. Node B acknowledges the FIN packet.
3. Node B sends a Finalize (FIN) packet
4. Node A acknowledges the FIN packet.

# Closing a Connection



**FIN WAIT 1**

**FIN WAIT 2**

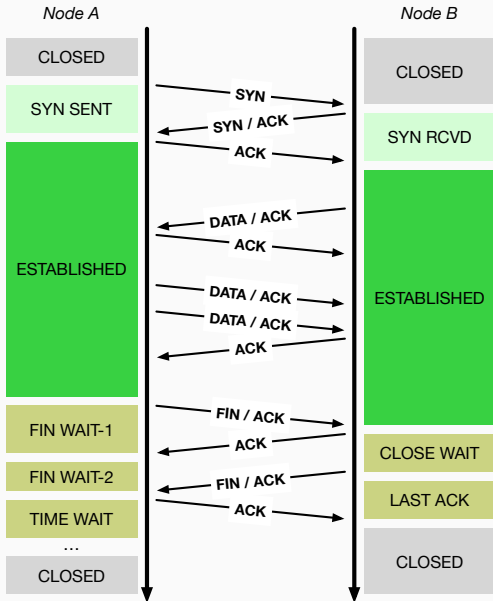
**TIME WAIT**

**CLOSE WAIT**

**LAST ACK** Awaiting client's final  
acknowledgment



# TCP State Machine



# DTrace and Networking Walkthrough

IPC Signals